

# RL-QEC: Harnessing Reinforcement Learning for Quantum Error Correction Advancements

Veeresh K  
Indian Institute of Science,  
Bangalore 560012, India  
veereshk@iisc.ac.in

Deepak S  
Cochin University of Science and Technology,  
India

Srinivas T  
Indian Institute of Science,  
Bangalore 560012, India

**Abstract-** With a particular emphasis on fixing fault defects caused by bitflip and depolarizing noise, this study explores the use of RL methods within the framework of quantum error correction. This work investigates the flexibility and independence of RL-trained agents in reducing bitflip mistakes by methodically looking at error rates and qubit lifetimes. Furthermore, we study how the top-performing agents are affected by depolarizing noise and present a strategic "Decoded Surface Code" method for fixing mistakes. The results show how RL works in handling different types of errors on its own, revealing subtle dynamics in the quantum system. Quantum error correction is a rapidly developing area, and this study adds to it by shedding light on how reinforcement learning could improve quantum computing's fault tolerance. The study's findings provide a basis for future research into more robust quantum computing systems and have ramifications for the creation of autonomous error-correcting algorithms.

**Keywords Used-** Quantum error correction, Reinforcement learning, Fault error mitigation, Bitflip noise, Depolarizing noise

## I. INTRODUCTION

Storage and handling of quantum information must be robust enough to resist mistakes introduced by physical qubits interacting with a noisy environment if large-scale quantum computing is to be implemented. One known approach to achieve this kind of robustness is by encapsulating the states of multiple physical qubits into a single logical qubit using a quantum error correcting code. Later on, this technique can be used to proactively find and fix errors [1, 2]. Although other quantum error-correcting codes exist, the topological quantum codes that are both feasible from an empirical standpoint [1-8] and that utilise solely local operations for error detection and correction [9-15] are of special interest. The surface code is an attractive option for large-scale fault-tolerant quantum computation because of its minimal overhead, locality requirements, and readily available ways to implement all the necessary logical gates [16, 17]. As a matter of fact, existing road maps recognise surface code-based approaches as the best bet for achieving strong quantum computing [18].




Fig. 1. Quantum error correction

Their lack of self-correction, the surface code, and the other realistic topological quantum error correcting codes are vulnerable to thermal noise that occurs naturally. Decoding algorithms are essential in any code-based approach to fault-tolerant quantum computation because of the need to actively detect and fix mistakes. The results of syndrome measurements, which diagnose problems on the physical qubits, are input into these algorithms, and they produce, as an output, a recommendation on how to fix any computation mistakes. In reality, these decoding algorithms need to be lightning-fast; more specifically, they need to be able to decode at a rate greater than the pace at which mistakes happen. Therefore, to understand quantum computing in practical settings, the development of decoding algorithms is crucial, as they are a major roadblock to the creation of fault-tolerant quantum computers. Note that the necessary syndrome measurements in any physically plausible scenario are likewise inherently inaccurate since they are derived from tiny quantum circuits. This is why, even if ideal syndrome measurements provide a great platform for developing decoding algorithms, any algorithm that wants to be useful in experiments needs to be able to handle these kinds of flawed observations. I must emphasise, nevertheless, that going from flawless syndrome decoders to full fault-tolerant decoders is no easy feat. Furthermore, these algorithms need to be able to handle experimentally realistic noise models and do calculations quickly enough to not be slowed down by the system's size (code distance) increasing.

Currently, the most difficult component is coming up with decoders that work in a fault-tolerant environment. Decoding algorithms are crucial for fault-tolerant quantum computation; hence many different techniques have been devised, all with the goal of meeting the empirically needed conditions. Although other methods based on techniques like the renormalization group and locally operated cellular automata [17-19] have also been proposed, algorithms based on minimum-weight perfect matching subroutines [19] are perhaps the most prevalent. These algorithms do a good job of solving the problem in theory, and they might be fast enough in the future with the help of parallelization techniques [19] or highly optimised hardware architectures [18], but for now, they're hard to put into practice, so new approaches are welcome.

An increasing number of areas within quantum physics are utilising machine learning methods. This includes efficiently representing many-body quantum states [12-17], identifying phase transitions, and autonomously designing new experimental setups. To tackle the issue of quick decoding, various decoders based on neural networks have been proposed. Due to the simplicity of inference techniques and the possibility of utilising dedicated hardware, neural network

decoders that utilise these specific chips may offer decoding speeds that are quick enough. Another issue with scalability is the difficulty of the training stage. Some novel methods, however, have the potential to alleviate this problem and facilitate the training of decoders for codes with code distances as high as 64. Finally, different proposals have demonstrated varying degrees of flexibility regarding the code and noise model. Despite the fact that the majority of neural network decoders only function in the idealised setting of syndrome measurements, there have been proposals that can manage circuit-level noise models. In fault-tolerant quantum memory, where the decoder has access to a final read-out of all physical qubits, references concentrate on this topic. Since logical gates, rather than measurements, follow decoding cycles in a computational setting, this method is not particularly applicable there. The absence of an algorithm or technique that satisfactorily satisfies all the essential experimental requirements, including speed, scalability, and accuracy, suggests that there is room for improvement, particularly in the ideal completely fault-tolerant computational environment. This is so even though there have been several decoding algorithms suggested and methods based on machine learning have shown great promise.

## II. METHODS AND SYSTEMATIC LAYOUT

### A. Earlier approach

An innovation known as a quantum repeater has made long-distance quantum communication a reality. At now, most protocols for quantum repeaters have been built with predefined quantum codes or graph states. The suggested general architecture for all-photon one-way quantum repeaters in this study is based on measurement-based error correction. All known Calderbank-Shor-Steane codes, including the recently discovered quantum low-density parity check (QLDPC) codes, can be modeled using this paradigm. In our decoding approach, the destination does the error correction using the collected data from the network measurements. Not only does this method simplify the local quantum operations at repeaters, but it also outperforms the standard protocols with independent repeaters. We statistically demonstrate, for instance, that the [6, 8] generalised bicycle code (as a compact but efficient QLDPC code) achieves comparable performance while substantially lowering resources. with the uses of by protecting logical qubits through designed dissipation, autonomous quantum error correction (AQEC) eliminates the need for frequent and error-prone measurement-feedback loops. The controllability and flexibility of bosonic code spaces make them attractive AQEC candidates, as single-photon loss is their major source of error. Although AQEC with bosonic code spaces has been shown to be feasible in principle in previous suggestions, these schemes usually rely on applying the Knill-Laflamme requirements exactly, which means that they necessitate the realisation of Hamiltonian distances  $d \geq 2$ . These methods are experimentally tough to implement because they involve many nonlinear interactions and control fields, which are equivalent to Hamiltonian distances.

### B. Current approach

A systematic approach is taken to improve the efficacy and autonomy of error correction protocols in quantum computing systems through the concept of using reinforcement learning (RL) to advance quantum error correction (QEC). It all starts with establishing a thorough knowledge of the quantum error

landscape, which involves naming important error types and the probability associated with them in the quantum system. Then, a framework for reinforcement learning is built so that the system can optimise its error correction procedures according to real-time feedback and performance evaluations. To motivate the quantum system to learn and execute effective actions for mistake correction, a reward system must be developed. To allow the RL algorithm to adapt to the specific quantum environment, it is trained iteratively using simulations or real quantum hardware. Extensive testing scenarios further demonstrate the robustness and generalizability of the RL model, ensuring its applicability across diverse quantum computing systems. By applying the flexibility and learning power of reinforcement learning to the problem of quantum error correction, this approach offers a potential way to improve the accuracy and efficiency of quantum computers.

A methodical strategy to strengthen quantum computing system resilience is to combine quantum error correction with reinforcement learning (RL) shown in figure 2. A thorough grasp of the error scenarios common to quantum systems is initially gained by performing a quantum error landscape study. A probabilistic evaluation of errors is a part of this analysis, with the goal of identifying and ranking the most important kinds of errors. The sorts of errors that have been found will then be used to create a specific RL algorithm for quantum error correction.

A meticulously defined action space, outlining the set of feasible activities the quantum system might perform for mistake correction, is a defining feature of the RL algorithm. The system can adjust on its own using real-time feedback from its quantum environment, thanks to the algorithm's continuous learning mechanisms. Through a combination of real-time feedback and continuous learning, a dynamic framework is created. This framework enables the RL algorithm to adapt its tactics for mistake correction as iterations progress. With the RL algorithm in practice, one must train it on both quantum simulators and real quantum hardware. Implementing quantum hardware guarantees real-world application and validation of the learned error correction rules, while simulated environments provide a controlled setting for algorithm refinement. Because quantum systems are inherently unpredictable, it is critical to analyse variability in quantum environments. The algorithm's robustness to changes in input parameters and external factors is evaluated by a sensitivity study. Examining the RL algorithm's output response allows us to gauge how well it decreases error rates and strengthens the quantum computing system as a whole. As the complexity of quantum systems and the number of qubits increases, it becomes increasingly important to examine the scalability of RL approaches used for quantum error correction. To further understand the impact and practicality of RL-based error correction in the quantum setups, it is important to study the computing overhead and integration difficulties of doing so. Dynamic changes in quantum error patterns and noise characteristics should also be able to be accommodated by scalable RL algorithms. In order to make RL-based techniques useful and effective in reducing mistakes in increasingly complicated quantum computing settings, it is crucial to solve these scaling issues.

### C. Traditional quantum error correction techniques versus the RL approach

In quantum error correction, traditional methods like surface codes, stabilizer codes, and quantum codes are well-established but often require many qubits and complex optimization. Reinforcement Learning (RL) offers a new approach, adapting strategies to hardware and noise models, potentially improving scalability, especially with limited qubits. Challenges include training complexity and generalization across diverse hardware. Integrating RL with traditional methods could enhance error correction. This analysis highlights the importance of understanding RL's advantages and limitations in quantum error correction, offering insights for future research.




Fig. 2. Methodological approach of current analysis

### III. PSEUDO ANALYSIS

**Note-** This pseudo code outlines the main processes of combining RL with quantum error correction, including analysing the quantum error landscape, identifying probabilistic errors, developing RL algorithms, implementing quantum hardware, analysing variability and sensitivity, and finally, the output response.

#### # Quantum Error Landscape Analysis

- `error_landscape = analyze_quantum_error()`

#### # Probabilistic Error Identification

- `error_probabilities = probabilistic_error_identification(error_landscape)`

#### # Reinforcement Learning Algorithm

- `rl_algorithm = RLAlgorithm()`
- `action_space = define_action_space()`
- `rl_algorithm.train(action_space, continuous_learning=True, real_time_feedback=True)`

#### # Quantum Hardware Implementation

- `quantum_hardware = QuantumHardware()`
- `simulation_data = simulate_quantum_system()`
- `real_data = quantum_hardware.run()`

#### # Variability Analysis

- `quantum_variability = analyze_variability(simulation_data, real_data)`

#### # Sensitivity Analysis

- `sensitivity_analysis = analyze_sensitivity(rl_algorithm, quantum_variability)`

#### # Output Response

- `output_response = rl_algorithm.execute(error_landscape)`

# End

### IV. IMPLEMENTATION AND RESULT LAYOUT




Fig. 3. Bitflip noise consideration

A quantum computing system is thoroughly examined in the given sequence in Figure 3, which shows an iterative or temporal progression denoted by a decreasing numeral sequence ( $10^5$ ,  $10^4$ ,  $10^3$ , etc.) and parts delineated with symbols. The error rates or probabilities represented by the numerical values of "Single Faulty Qubit" and "Average Qubit Lifetime" (0.002, 0.004, etc.) provide light on the dependability of individual qubits and the lifespan of the system as a whole. Keeping the physical and measured probabilities equal is crucial for precise quantum computing, as the notation "Pphys = Pmeas" emphasises. Over the course of the experiment, the error rates or probabilities may fluctuate, as seen by the observed decreasing trend, subsequent increase, and stabilisation. Assigning numerical values to "Agent trained at" and "Bitflip Noise - All Agents" implies incorporating a reinforcement learning agent that has been trained under bitflip noise conditions, with these values representing performance metrics or outcomes that have been derived from this training. This detailed technical analysis reveals dynamic features of the quantum system's behaviour by combining observed probabilities with the impact of bitflip noise on an RL agent that has been trained. It would be helpful to have more background information in order to fully grasp the particular experimental setups and aims of this investigation.






Fig. 4. Depolarising noise consideration

A quantum computing system is analysed in the sequence shown in Figure 4, with numerical values and labels. Declining numerical progression ( $10^5$ ,  $10^4$ ,  $10^3$ , etc.) denotes progressive inquiry, maybe temporal or iterative. The "Single Faulty Qubit" and "Average Qubit Lifetime" values, as well as " $P_{phys} = P_{meas}$ ," emphasises qubit reliability and quantum computations' alignment between physical and measured probabilities. "(b) Depolarising Noise All Agents" suggests a thorough system resilience evaluation by considering depolarizing noise affecting all agents. A trained agent is involved when "Agent trained at" precedes numeric numbers, which may represent performance metrics or outcomes under certain noise settings. This sequence analyses error rates, qubit lifetime, and depolarizing noise's effect on a trained reinforcement learning agent in a quantum computing system.

The given chain of events expands the technical study of the quantum computer system by presenting in Figure 5 a new setting called "(c) Bitflip Noise - Best Performing Agents" and then making reference to a "Decoded Surface Code." A continual inquiry or experimental iteration is suggested by the numeral sequence that descends ( $10^5$ ,  $10^4$ ,  $10^3$ , etc.).




Fig. 5. : Decoded surface code for Bitflip noise

There has been a consistent emphasis on error rates and qubit lifetime, as seen by the figures associated with "Single Faulty Qubit" and "Average Qubit Lifetime" (0.002, 0.004, etc.). Aligning physical and measured probabilities is crucial in quantum computations, as the formula " $P_{phys} = P_{meas}$ " emphasises once again. The inclusion of "(c) Bitflip Noise - Best Performing Agents" suggests a focused assessment of system robustness and a detailed examination of how bitflip noise affects the top-performing agents. The following discussion of "Decoded Surface Code" presents a possible strategy for dealing with quantum mistakes. Using a decoded surface code, bitflip noise's effect on top-performing agents,

qubit lifetime, and error characteristics are all covered in this multi-faceted investigation.




Fig. 6. Decoded surface code for Depolarising noise

The stated sequence is shown in Figure 6 with "(d) Depolarizing Noise Best Performing Agents," and then a "Decoded Surface Code" is mentioned. The numerical sequence continues to fall, indicating a sequential exploration or experimental iteration ( $10^5$ ,  $10^4$ ,  $10^3$ , etc.). There is a continuous investigation into error rates and qubit lifetime as the values linked to "Single Faulty Qubit" and "Average Qubit Lifetime" (0.002, 0.004, etc.) remain. " $P_{phys} = P_{meas}$ " is a notation that stresses the significance of matching physical and measured probabilities in quantum computing. Incorporating "(d) Depolarizing Noise Best Performing Agents" into the discussion implies a targeted analysis of how depolarizing noise affects agents that do exceptionally well. The mention of a "Decoded Surface Code," which might be a method for fixing mistakes, further strengthens this idea. Under certain noise situations, the digits that go into the double digits (10+) might indicate the quantum system's scale or an enhanced effect. This brief examination highlights an extensive study of error properties, qubit lifetime, the effect of depolarizing noise on top-performing agents, and the use of a decoded surface code to reduce errors.

## V. CONCLUSION-

With a focus on bitflip and depolarizing noise, this study explored quantum error correction via reinforcement learning. Agents taught with reinforcement learning were able to independently reduce bitflip noise-induced error rates by systematically analysing qubit lifetimes and error rates. In addition to presenting the idea of a "Decoded Surface Code" as a strategy for error correction, the study went on to examine how depolarizing noise affected the top-performing agents. The complex dynamics of the quantum system under certain noise conditions were highlighted by the trends in numerical values, particularly the rise into the double digits when depolarizing noise was present. Taken as a whole, these results highlight the potential of reinforcement learning to improve quantum computing's fault tolerance and pave the way for autonomous error correction algorithms. Future research endeavours can build upon this study to refine and expand upon these error correction approaches, with the goal of developing more resilient and dependable quantum computing systems. This is especially important as quantum technologies continue to improve.

## REFERENCES

- [1] Terhal, Barbara M. "Quantum error correction for quantum memories." *Reviews of Modern Physics* 87, no. 2 (2015): 307.

- [2] Campbell, Earl T., Barbara M. Terhal, and Christophe Vuillot. "Roads towards fault-tolerant universal quantum computation." *Nature* 549, no. 7671 (2017): 172-179.
- [3] Kitaev, A. Yu. "Fault-tolerant quantum computation by anyons." *Annals of physics* 303, no. 1 (2003): 2-30.
- [4] Dennis, Eric, Alexei Kitaev, Andrew Landahl, and John Preskill. "Topological quantum memory." *Journal of Mathematical Physics* 43, no. 9 (2002): 4452-4505.
- [5] Alexeev, Yuri, Dave Bacon, Kenneth R. Brown, Robert Calderbank, Lincoln D. Carr, Frederic T. Chong, Brian DeMarco et al. "Quantum computer systems for scientific discovery." *PRX Quantum* 2, no. 1 (2021): 017001.
- [6] Luckow, Andre, Johannes Klepsch, and Josef Pichlmeier. "Quantum computing: Towards industry reference problems." *Digitale Welt* 5 (2021): 38-45.
- [7] Rowell, Eric, and Zhenghan Wang. "Mathematics of topological quantum computing." *Bulletin of the American Mathematical Society* 55, no. 2 (2018): 183-238.
- [8] Nimbe, Peter, Benjamin Asubam Weyori, and Adebayo Felix Adekoya. "Models in quantum computing: a systematic review." *Quantum Information Processing* 20, no. 2 (2021): 80.
- [9] Sivak, V. V., A. Eickbusch, H. Liu, B. Royer, I. Tsoutsios, and M. H. Devoret. "Model-free quantum control with reinforcement learning." *Physical Review X* 12, no. 1 (2022): 011059.
- [10] Skolik, Andrea, Stefano Mangini, Thomas Bäck, Chiara Macchiavello, and Vedran Dunjko. "Robustness of quantum reinforcement learning under hardware errors." *EPJ Quantum Technology* 10, no. 1 (2023): 1-43.
- [11] Wang, Hao-Wen, Qian Cao, Yun-Jia Xue, Li Ding, Han-Yang Liu, Yu-Min Dong, and Hong-Yang Ma. "Determining quantum topological semion code decoder performance and error correction effectiveness with reinforcement learning." *Frontiers in Physics* 10 (2022): 981225.
- [12] Martín-Guerrero, José D., and Lucas Lamata. "Reinforcement learning and physics." *Applied Sciences* 11, no. 18 (2021): 8589.
- [13] Jerbi, Sofiene, Lea M. Trenkwalder, Hendrik Poulsen Nautrup, Hans J. Briegel, and Vedran Dunjko. "Quantum enhancements for deep reinforcement learning in large spaces." *PRX Quantum* 2, no. 1 (2021): 010328.
- [14] Wei, Qing, Hailan Ma, Chunlin Chen, and Daoyi Dong. "Deep reinforcement learning with quantum-inspired experience replay." *IEEE Transactions on Cybernetics* 52, no. 9 (2021): 9326-9338.
- [15] Chen, Samuel Yen-Chi, Chao-Han Huck Yang, Jun Qi, Pin-Yu Chen, Xiaoli Ma, and Hsi-Sheng Goan. "Variational quantum circuits for deep reinforcement learning." *IEEE Access* 8 (2020): 141007-141024.
- [16] Mackeprang, Jelena, Durga B. Rao Dasari, and Jörg Wrachtrup. "A reinforcement learning approach for quantum state engineering." *Quantum Machine Intelligence* 2 (2020): 1-14.
- [17] Borah, Sangkha, Bijita Sarma, Michael Kewming, Gerard J. Milburn, and Jason Twamley. "Measurement-based feedback quantum control with deep reinforcement learning for a double-well nonlinear potential." *Physical review letters* 127, no. 19 (2021): 190403.
- [18] Moro, Lorenzo, Matteo GA Paris, Marcello Restelli, and Enrico Prati. "Quantum compiling by deep reinforcement learning." *Communications Physics* 4, no. 1 (2021): 178.
- [19] Ostaszewski, Mateusz, Lea M. Trenkwalder, Wojciech Masarczyk, Eleanor Scerri, and Vedran Dunjko. "Reinforcement learning for optimization of variational quantum circuit architectures." *Advances in Neural Information Processing Systems* 34 (2021): 18182-18194.